

Retail Pricing & Demand Forecasting System

A machine learning pipeline for demand forecasting and dynamic pricing optimization for multi-store retail operations.

Overview

This system helps retailers optimize pricing strategies by:

- Forecasting demand for products across stores
- Optimizing prices to maximize revenue, units sold, or margins
- Simulating the impact of pricing strategies

Dataset

The system works with retail data containing:

- **5 stores** (S01-S05) and **3 SKUs** (SKU001-SKU003)
- Daily sales data with pricing, promotions, and external factors
- Pre-split train/test data for model validation

Quick Start

1. Train the model:

```
python pipeline.py train
```

1. Optimize prices for maximum revenue:

```
python pipeline.py optimize --objective revenue --out pricing_results.csv
```

1. Simulate a custom price plan:

```
python pipeline.py simulate --price-plan my_prices.csv --out  
simulation_results.csv
```

Project Structure

```
├── src/  
│   ├── data_loader.py    # Data loading utilities  
│   ├── features.py       # Feature engineering  
│   ├── models/  
│   │   ├── forecaster.py # Demand forecasting model  
│   │   └── pricing/  
│   │       └── optimizer.py # Price optimization engine  
├── pipeline.py           # Main CLI interface  
└── README.md
```

Model Performance

- **Algorithm:** Gradient Boosting Regressor
- **Performance:** ~22% MAPE on test set
- **Key Features:** Historical demand patterns, price discounts, promotional depth

Business Insights

- **Promotional Impact:** 65-115% demand increase with promotions
- **Price Elasticity:** SKU003 most sensitive (-0.71), SKU001 least (-0.48)
- **Store Performance:** Clear hierarchy with S05 outperforming S01 by 94%

Price Optimization Results

For SKU001 in Store S01:

- **Revenue Maximization:** 20% discount increases revenue by 60%
- **Margin Maximization:** 25% premium pricing maximizes per-unit profit

- **Trade-off:** Cannot optimize revenue and margin simultaneously

Requirements

pandas
scikit-learn
numpy
joblib

Usage Examples

Training with Custom Date Range

```
python pipeline.py train --train-start 2024-01-01 --train-end 2024-09-30
```

Multi-Objective Optimization

```
python pipeline.py optimize --objective margin --out margin_results.csv
```

Limitations

- Assumes fixed cost structure (\$5 per unit for calculations)
- 22% MAPE indicates room for improvement in forecast accuracy
- No inventory constraints or competitive response modeling
- Results based on historical patterns may not capture future market changes

Future Enhancements

- Incorporate inventory constraints
- Add competitor response modeling
- Implement hierarchical forecasting
- Include customer behavior dynamics

This implementation demonstrates core ML pipeline capabilities for retail pricing optimization while maintaining simplicity and interpretability.